EXPERIMENT NO. 03

Aim: To understand the Kubernetes Cluster Architecture, install and Spin Up a Kubernetes Cluster on Linux Machines/Cloud Platforms.

Aim: To install Kubectl and execute Kubectl commands to manage the Kubernetes cluster and deploy YourFirst Kubernetes

STEP 1:

Check security group,delete all SG only keep default

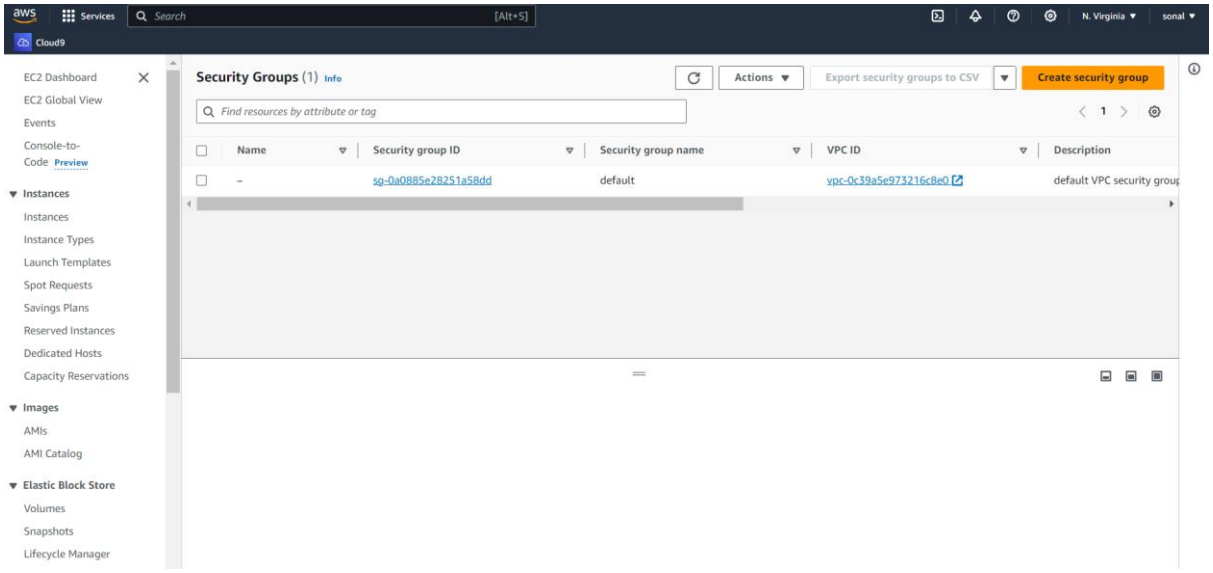

Create 2 instance

Create key pair

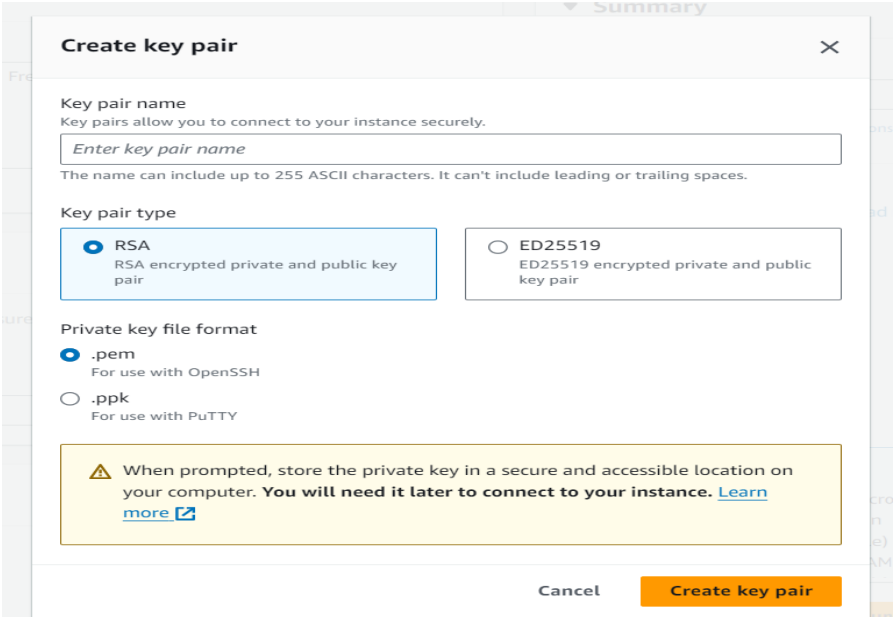

Create security group and allow traffic

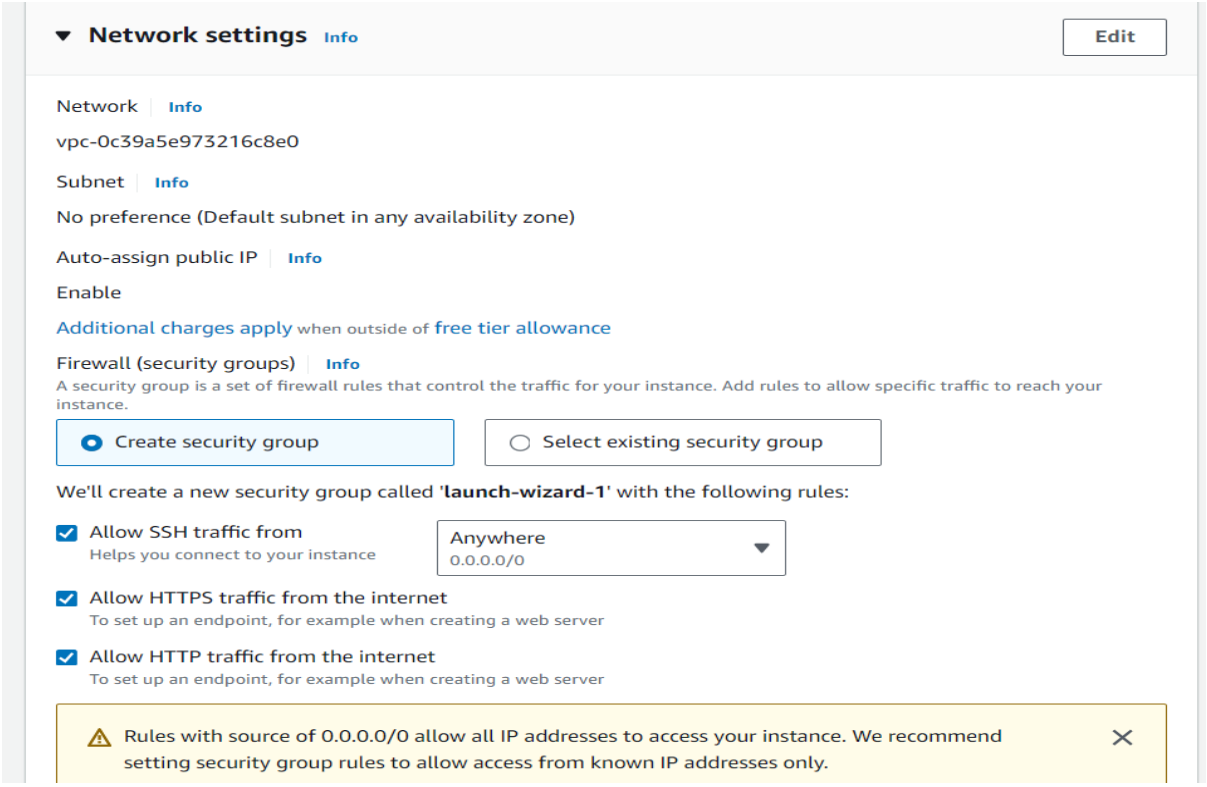

Launch Instance

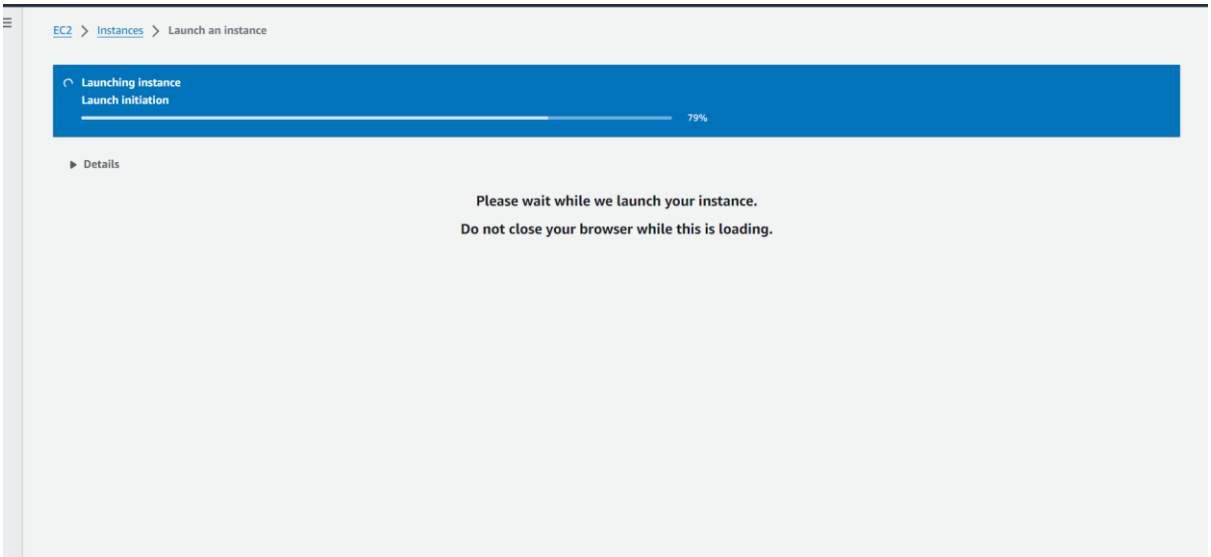

Check security group of both instances

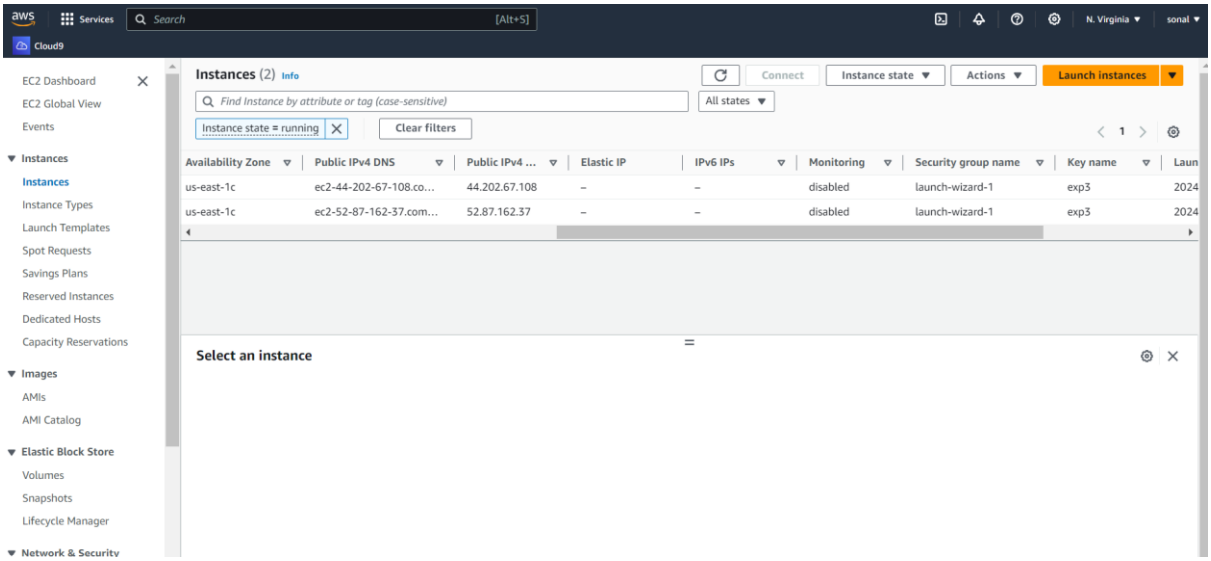

Here security group is launch wizard-1

Now go to security group from left pane

Click on Security group id of Launch wizard-1

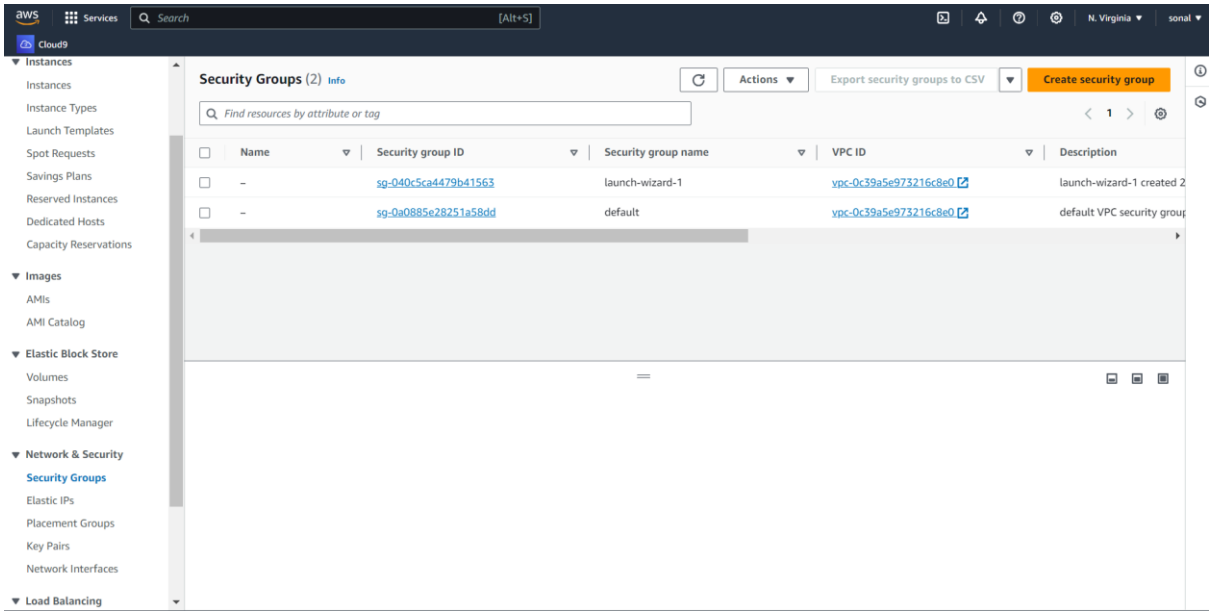

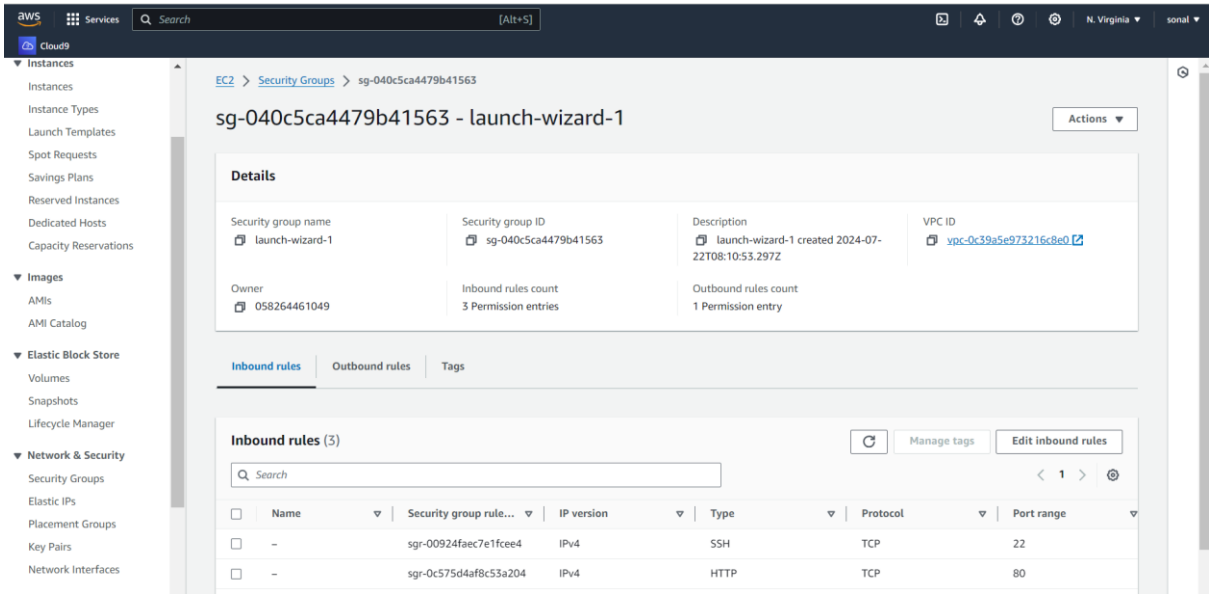

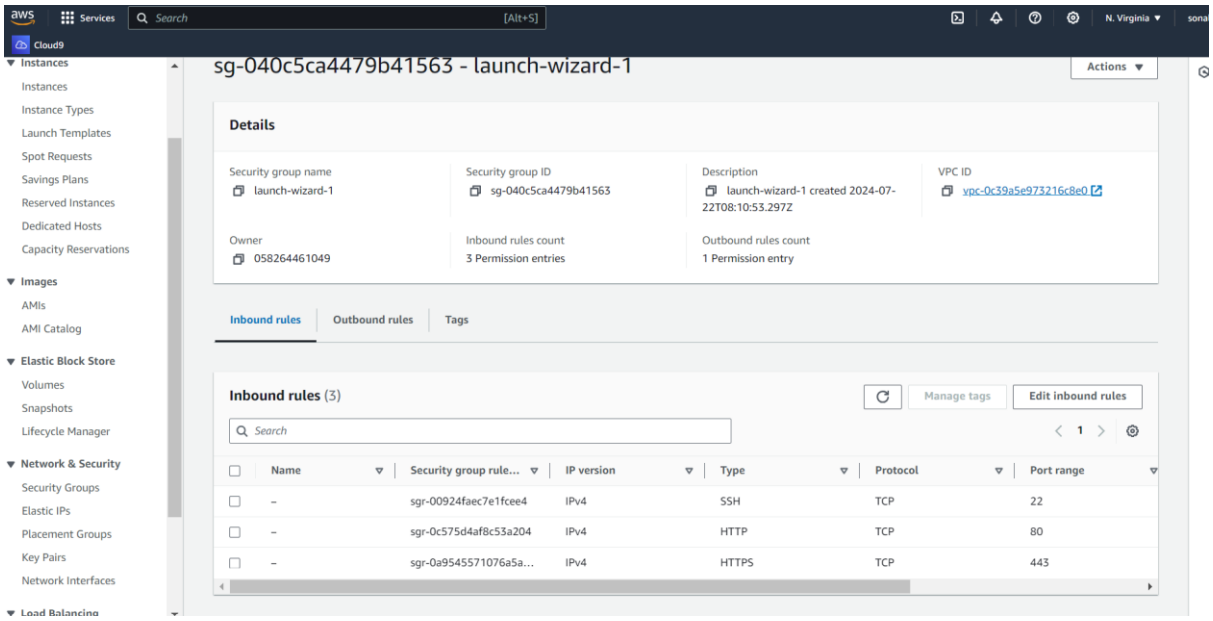

Edit inbound rule

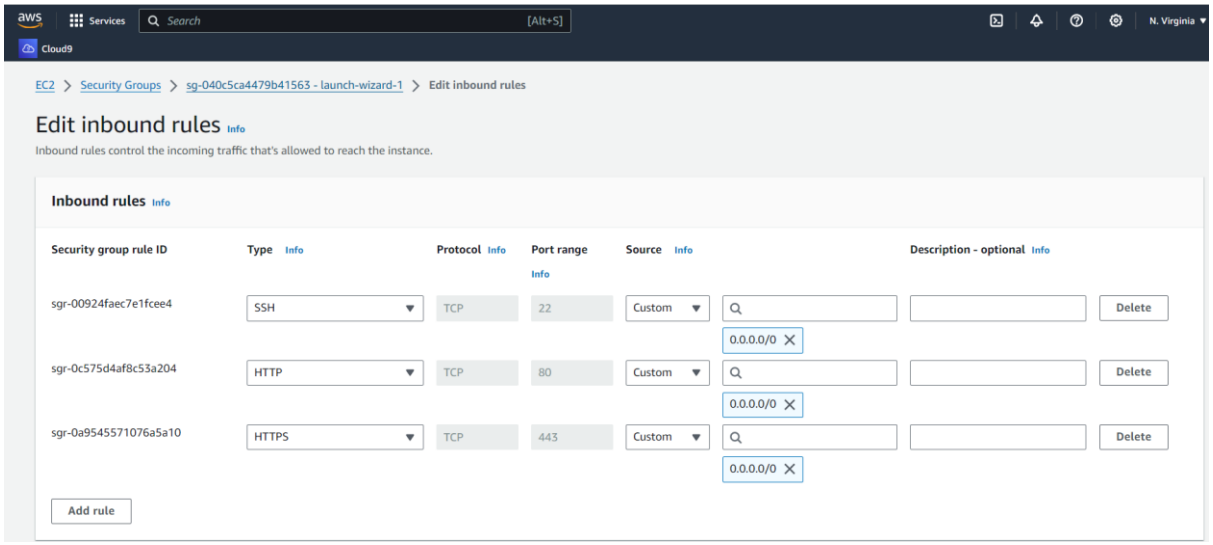

Delete all rules

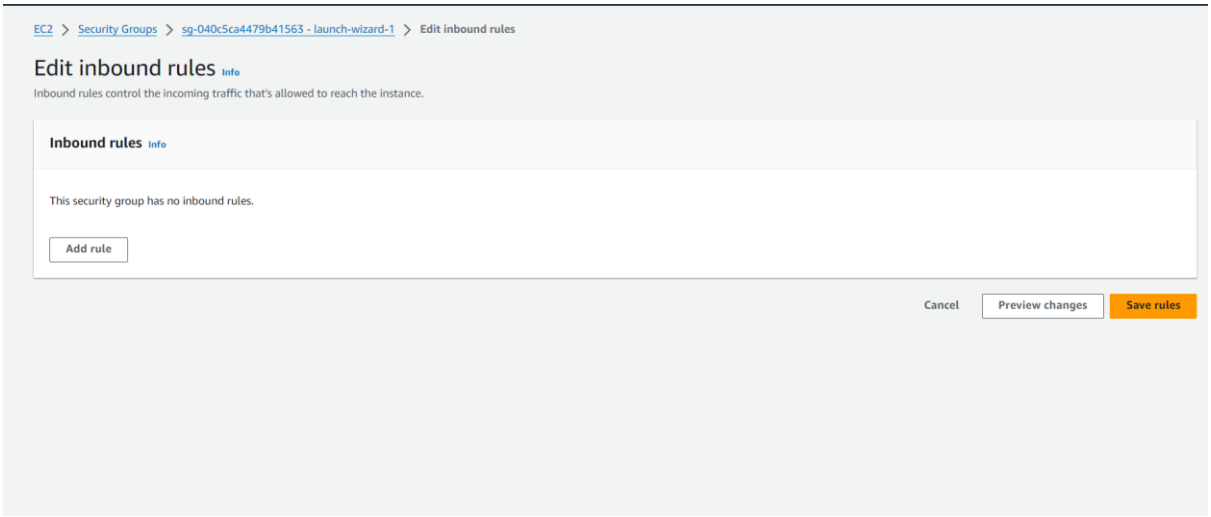

Add new rule

Select

ALL traffic

Anywhere IPV4

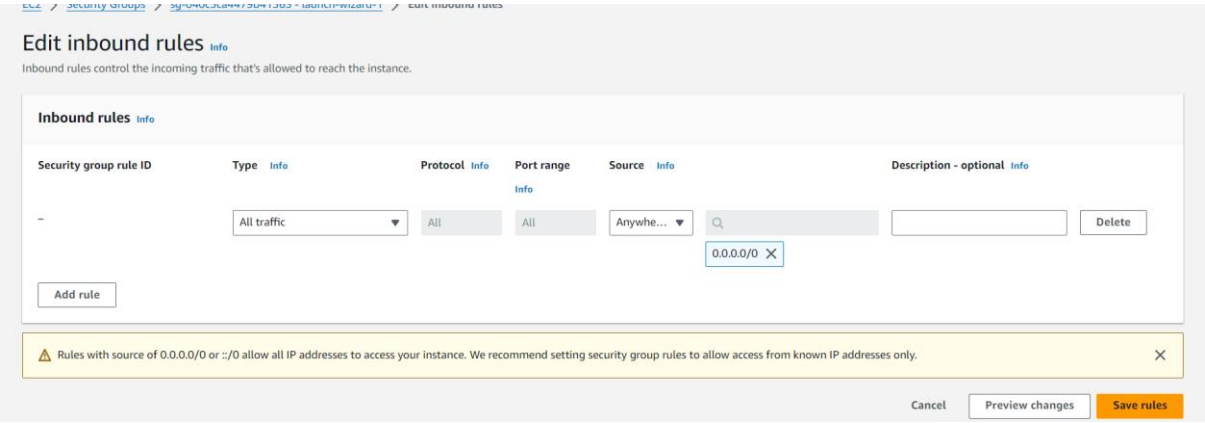

Now save rules

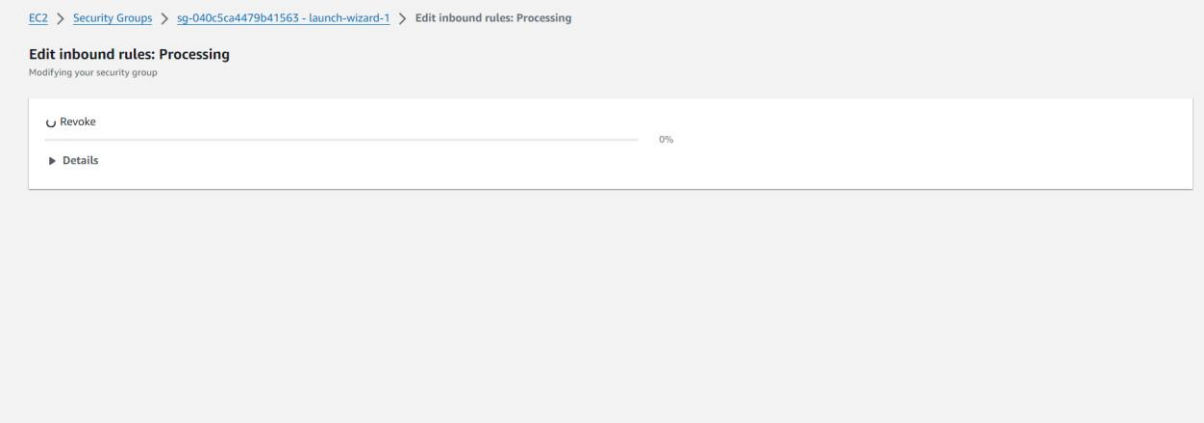

Name the instances

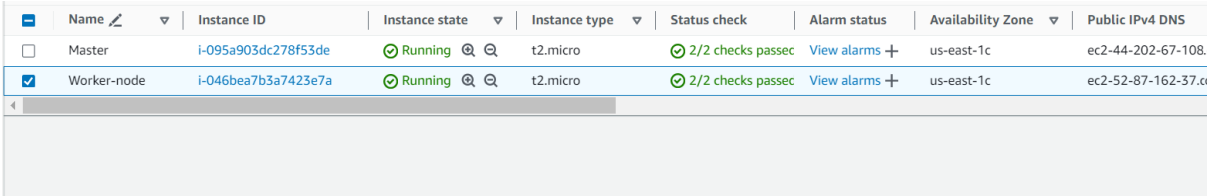

Select Master and connect

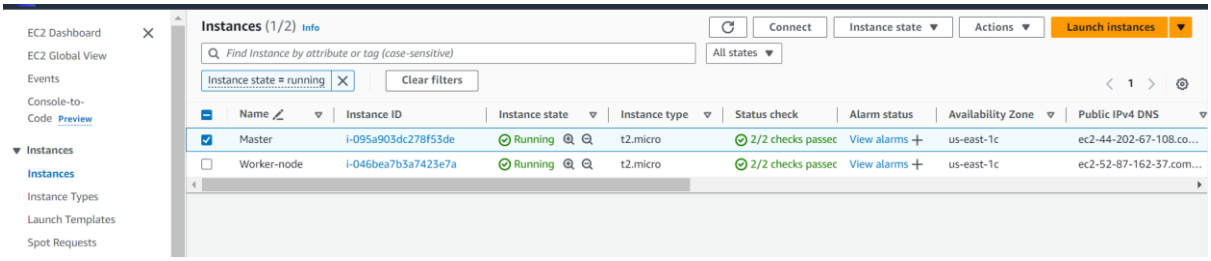

Click on connect

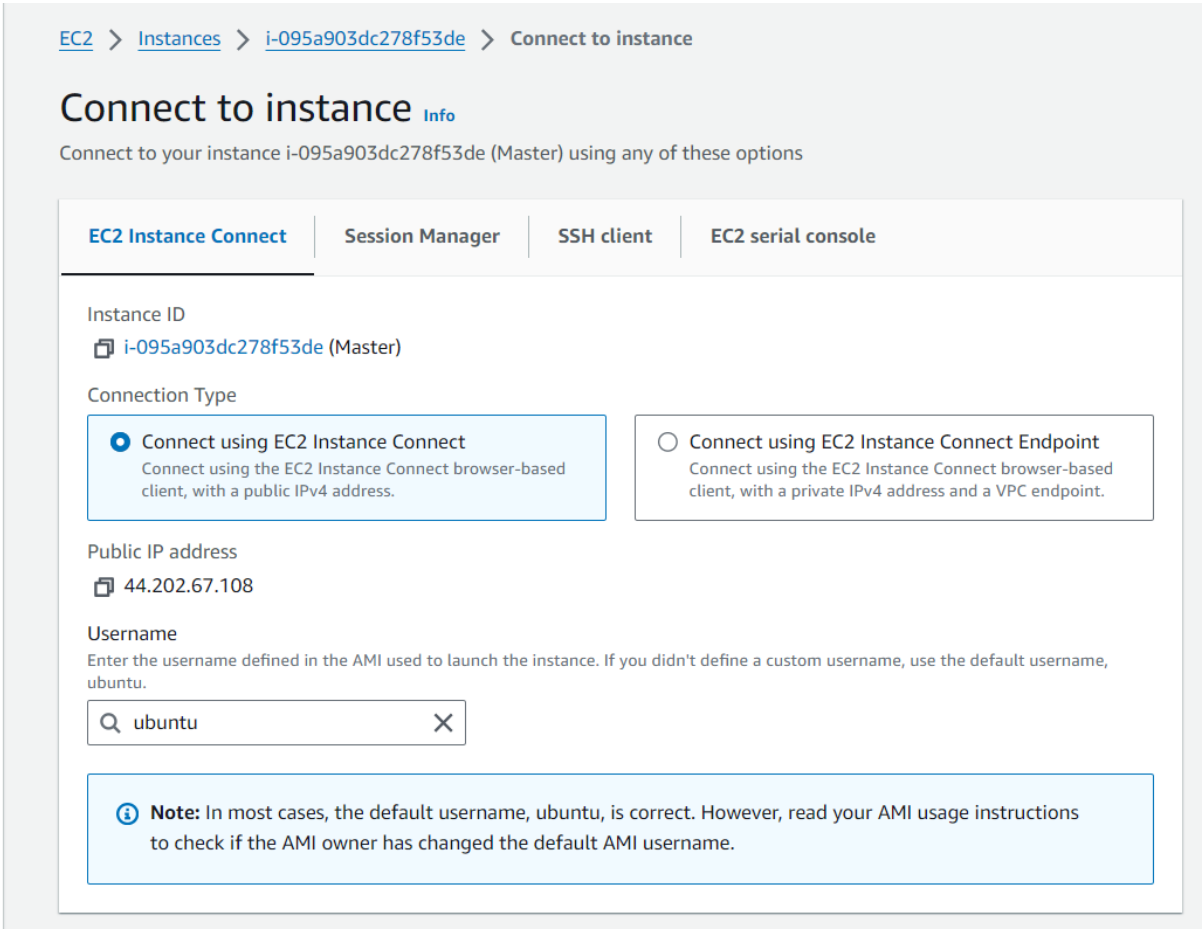

After Connecting

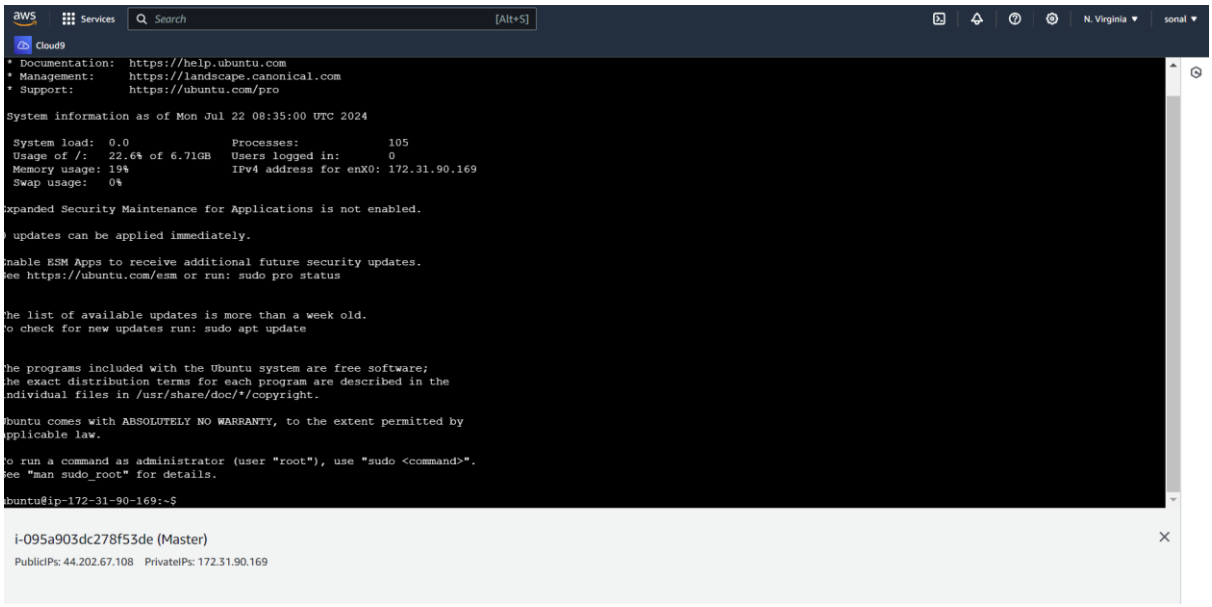

Sameway connect to worker-node

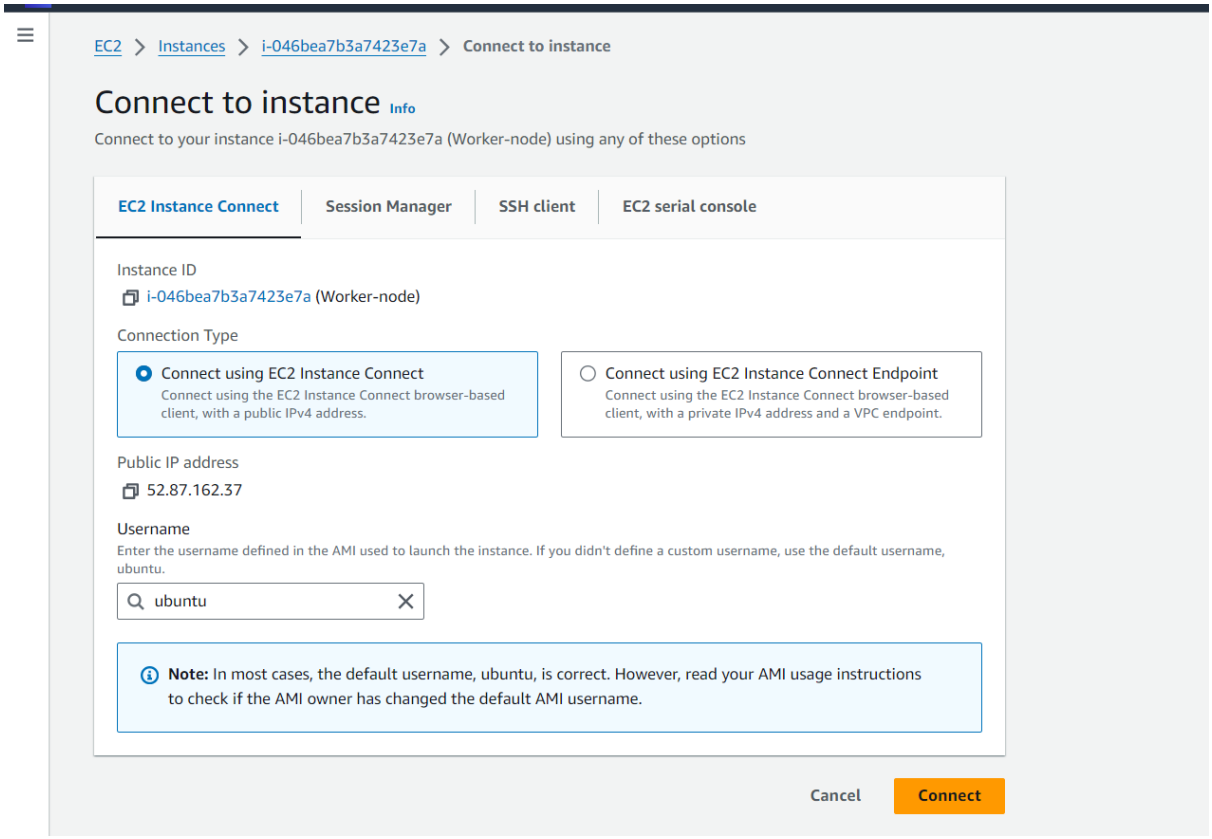

Step 2:

Assign Unique Hostname for Each Server Node

$ sudo hostnamectl set-hostname master-node

Than exit

Refresh

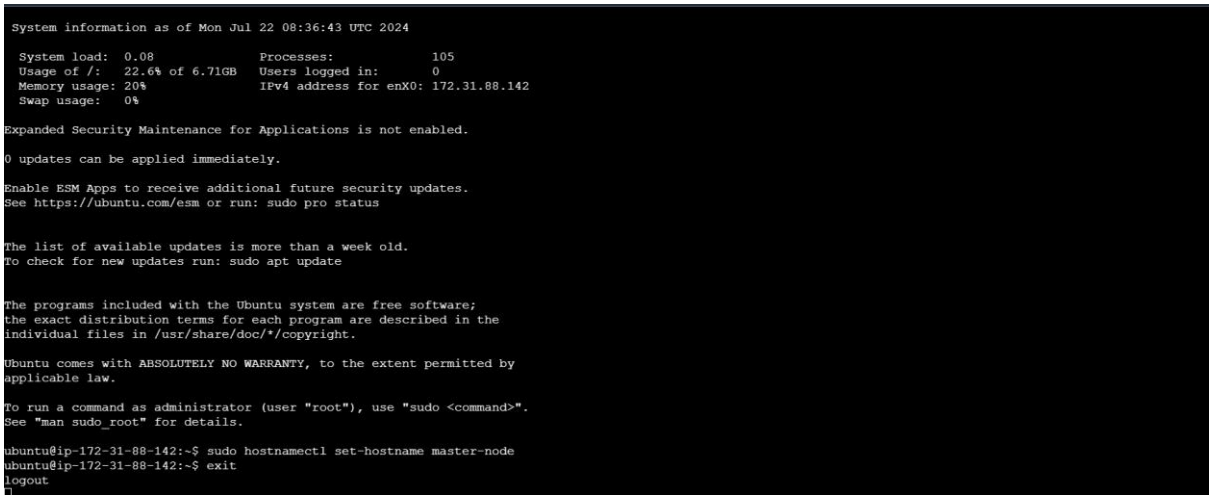

Next, set a worker node hostname by entering the following on the worker server:

$ sudo hostnamectl set-hostname worker1

STEP 3:

On both master and worker1

$ sudo apt-get update

STEP 4:

On both master and worker1

Install docker

sudo apt-get install docker.io

STEP 5 : Start and Enable Docker

Set Docker to launch at boot by entering the following:

$ sudo systemctl enable docker

$ sudo systemctl status docker

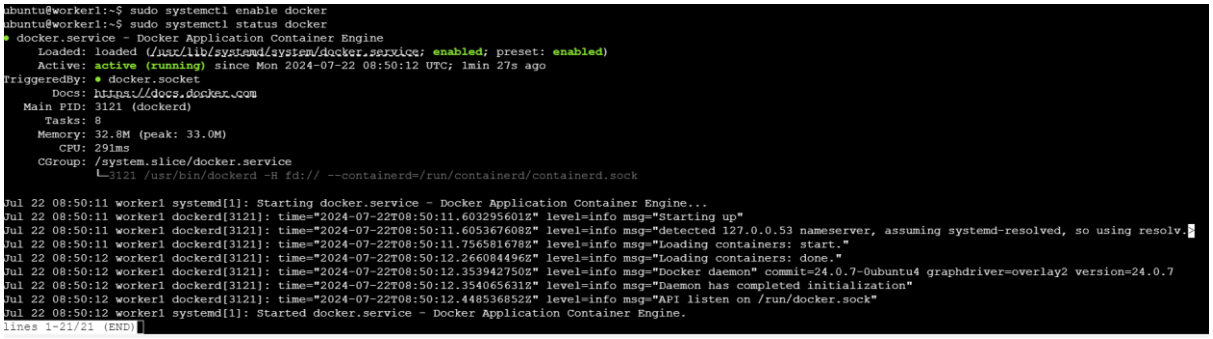

Ctrl+c

Clear

sudo systemctl start docker

STEP 6 Install Kubernetes

https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/install-kubeadm/

```
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
ca-certificates is already the newest version (20240203).
ca-certificates set to manually installed.
curl is already the newest version (8.5.0-2ubuntu10.1).
curl set to manually installed.
gpg is already the newest version (2.4.4-2ubuntu17).
gpg set to manually installed.
The following NEW packages will be installed:
  apt-transport-https
0 upgraded, 1 newly installed, 0 to remove and 22 not upgraded.
Need to get 3974 B of archives.
After this operation, 35.8 kB of additional disk space will be used.
Get:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/universe amd64 apt-transport-https all 2.7.14build2 [3974 B]
Fetched 3974 B in 0s (268 kB/s)
Selecting previously unselected package apt-transport-https.
(Reading database ... 68106 files and directories currently installed.)
Preparing to unpack .../apt-transport-https_2.7.14build2_all.deb ...
Unpacking apt-transport-https (2.7.14build2) ...
Setting up apt-transport-https (2.7.14build2) ...
Scanning processes...
Scanning linux images...

Running kernel seems to be up-to-date.

No services need to be restarted.

No containers need to be restarted.

No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on this host
```

Signining key

```
ubuntu@worker1:~$ curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.30/deb/Release.key | sudo gpg --dearmor -o /etc/apt/keyrings/kubernetes-apt-keyring.gpg
ubuntu@worker1:~$
```

Cloud9

```
Preparing to unpack .../5-kubernetes-cni_1.4.0-1.1_amd64.deb ...
Unpacking kubernetes-cni (1.4.0-1.1) ...
Selecting previously unselected package socat.
Preparing to unpack .../6-socat_1.8.0.0-4build3_amd64.deb ...
Unpacking socat (1.8.0.0-4build3) ...
Selecting previously unselected package kubelet.
Preparing to unpack .../7-kubelet_1.30.3-1.1_amd64.deb ...
Unpacking kubelet (1.30.3-1.1) ...
Setting up conntrack (1:1.4.8-1ubuntu1) ...
Setting up kubectl (1.30.3-1.1) ...
Setting up ebtables (2.0.11-6build1) ...
Setting up socat (1.8.0.0-4build3) ...
Setting up cri-tools (1.30.0-1.1) ...
Setting up kubernetes-cni (1.4.0-1.1) ...
Setting up kubeadm (1.30.3-1.1) ...
Setting up kubelet (1.30.3-1.1) ...
Processing triggers for man-db (2.12.0-4build2) ...
Scanning processes...
Scanning linux images...

Running kernel seems to be up-to-date.

No services need to be restarted.

No containers need to be restarted.

No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on this host.
ubuntu@worker1:~$ sudo apt-mark hold kubelet kubeadm kubectl
kubelet set on hold.
kubeadm set on hold.
kubectl set on hold.
```

Step 7 : Begin Kubernetes Deployment

$ sudo swapoff –a

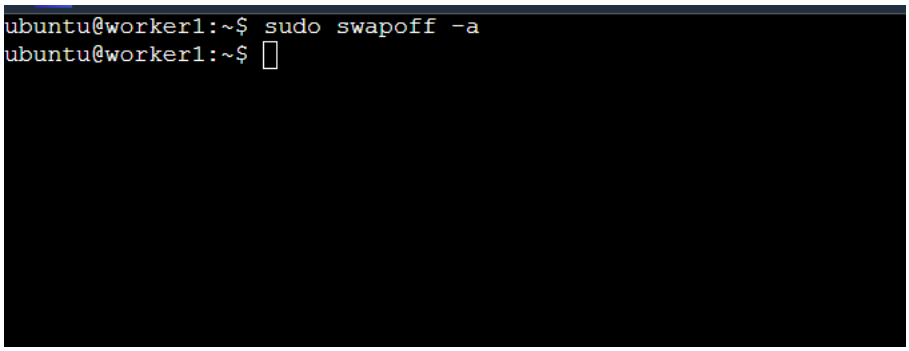

STEP 8:

Initialize Kubernetes on Master Node

$ sudo kubeadm init --pod-network-cidr=10.244.0.0/16 --ignore-preflight-errors=all

If getting error

Run below code on both

If the kubeadm init command ran without error then ignore this part. If you receive this error "kubelet isn't running or healthy", then do the following.
Create file daemon.json in /etc/docker/ and add following lines in the file.

```
{
   "exec-opts": ["native.cgroupdriver=systemd"]
}
```

And run the following commands.
Do this on both master and worker nodes.

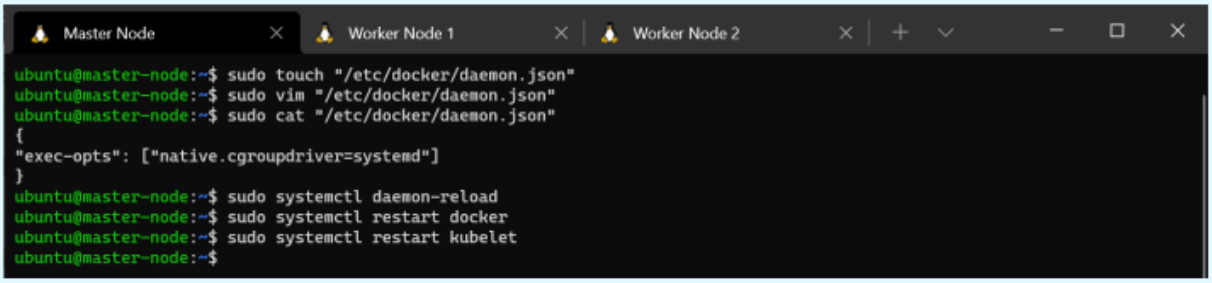

After this run sudo kubeadm reset command and then the init or join command.

$ sudo touch "/etc/docker/daemon.json"

$ sudo nano "/etc/docker/daemon.json"

ADD code

CTRL+ O enter CTRL X

$ sudo cat "/etc/docker/daemon.json"

{

"exec-opts": ["native.cgroupdriver=systemd"]

}

$ sudo systemctl daemon-reload

$ sudo systemctl restart docker

$ sudo systemctl restart kubelet

$ sudo kubeadm reset

STEP 9 on master node

sudo kubeadm init --pod-network-cidr=10.244.0.0/16 --ignore-preflight-errors=all

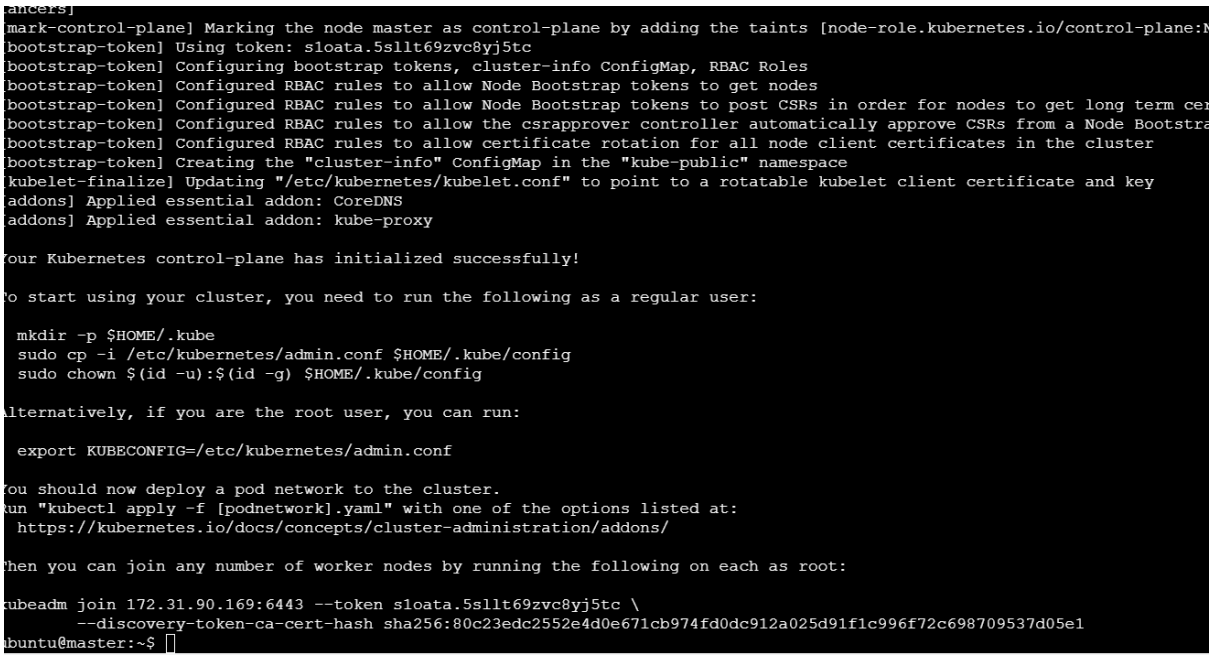

Next, enter the following to create a directory for the cluster: (Master)

kubernetes-master $ mkdir -p $HOME/.kube

kubernetes-master $ sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config

kubernetes-master $ sudo chown $(id -u):$(id -g) $HOME/.kube/config

copy

**kubeadm join 172.31.90.169:6443 --token s1oata.5sllt69zvc8yj5tc **

--discovery-token-ca-cert-hash sha256:80c23edc2552e4d0e671cb974fd0dc912a025d91f1c996f72c698709537d05e1

Makeit as and copy in worker after flannel is created on master(after step 10)

kubeadm join 172.31.90.169:6443 --token s1oata.5sllt69zvc8yj5tc --discovery-token-ca-cert-hash sha256:80c23edc2552e4d0e671cb974fd0dc912a025d91f1c996f72c698709537d05e1

it will give error

sudo kubeadm join 172.31.90.169:6443 --token s1oata.5sllt69zvc8yj5tc --discovery-token-ca-cert-hash sha256:80c23edc2552e4d0e671cb974fd0dc912a025d91f1c996f72c698709537d05e1 --ignore-preflight-errors=all

---

STEP 10 Copy weblink from masternode

https://kubernetes.io/docs/concepts/cluster-administration/addons/

goto flannel

copy this command and paste on master

For Kubernetes v1.17+

Deploying Flannel with kubectl

```
kubectl apply -f https://github.com/flannel-io/flannel/releases/latest/download/kube-flannel.yml
```

If you use custom `podCIDR` (not `10.244.0.0/16`) you first need to download the above manifest and modify the network to match your one.

Flannel created

```
Then you can join any number of worker nodes by running the following on each as root:

kubeadm join 172.31.90.169:6443 --token s1oata.5sllt69zvc8yj5tc \
        --discovery-token-ca-cert-hash sha256:80c23edc2552e4d0e671cb974fd0dc912a025d91f1c996f72c698709537d05e1
ubuntu@master:~$ mkdir -p $HOME/.kube
ubuntu@master:~$ sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
ubuntu@master:~$ sudo chown $(id -u):$(id -g) $HOME/.kube/config
ubuntu@master:~$ ^C
ubuntu@master:~$ ^C
ubuntu@master:~$ kubectl apply -f https://github.com/flannel-io/flannel/releases/latest/download/kube-flannel.yml
namespace/kube-flannel created
serviceaccount/flannel created
clusterrole.rbac.authorization.k8s.io/flannel created
clusterrolebinding.rbac.authorization.k8s.io/flannel created
configmap/kube-flannel-cfg created
daemonset.apps/kube-flannel-ds created
ubuntu@master:~$
```

Step 11:

AFTER

```
r endpoint \"unix:///var/run/containerd/containerd.sock\": rpc error: code = Unavailable desc = connection error: desc = \"transport: Error while dialing: dia
/run/containerd/containerd.sock: connect: permission denied\""
, error: exit status 1
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
error execution phase kubelet-start: couldn't save bootstrap-kubelet.conf to disk: open /etc/kubernetes/bootstrap-kubelet.conf: permission denied
To see the stack trace of this error execute with --v=5 or higher
ubuntu@worker1:~$ sudo ^[[200~kubeadm join 172.31.90.169:6443 --token s1oata.5sllt69zvc8yj5tc --discovery-token-ca-cert-hash sha256:80c23edc2552e4d0e671cb974f
d91f1c996f72c698709537d05e1 --ignore-preflight-errors=all
sudo: kubeadm: command not found
ubuntu@worker1:~$ sudo kubeadm join 172.31.90.169:6443 --token s1oata.5sllt69zvc8yj5tc --discovery-token-ca-cert-hash sha256:80c23edc2552e4d0e671cb974fd0dc912
96f72c698709537d05e1 --ignore-preflight-errors=all
[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-check] Waiting for a healthy kubelet. This can take up to 4m0s
[kubelet-check] The kubelet is healthy after 1.004306297s
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap

This node has joined the cluster:
* Certificate signing request was sent to apiserver and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.

ubuntu@worker1:~$
```

Step 12:

Reboot both instances

```
Last login: Mon Jul 22 08:44:23 2024 from 18.206.107.29
ubuntu@master:~$ kubectl get pods --all-namespaces
NAMESPACE      NAME                              READY   STATUS             RESTARTS       AGE
kube-flannel   kube-flannel-ds-qtfbw             1/1     Running            2 (66s ago)    16m
kube-flannel   kube-flannel-ds-t9v2g             1/1     Running            2 (65s ago)    6m29s
kube-system    coredns-7db6d8ff4d-2h5b4          1/1     Running            1 (119s ago)   30m
kube-system    coredns-7db6d8ff4d-sfzs8          1/1     Running            1 (119s ago)   30m
kube-system    etcd-master                       1/1     Running            1 (119s ago)   31m
kube-system    kube-apiserver-master             1/1     Running            1 (119s ago)   31m
kube-system    kube-controller-manager-master    1/1     Running            1 (119s ago)   31m
kube-system    kube-proxy-8kd97                  0/1     CrashLoopBackOff   5 (3s ago)     6m29s
kube-system    kube-proxy-9x78m                  0/1     CrashLoopBackOff   10 (21s ago)   30m
kube-system    kube-scheduler-master             1/1     Running            1 (119s ago)   31m
ubuntu@master:~$ kubectl get pods --all-namespaces
NAMESPACE      NAME                              READY   STATUS    RESTARTS        AGE
kube-flannel   kube-flannel-ds-qtfbw             1/1     Running   2 (118s ago)    17m
kube-flannel   kube-flannel-ds-t9v2g             1/1     Running   2 (117s ago)    7m21s
kube-system    coredns-7db6d8ff4d-2h5b4          1/1     Running   1 (2m51s ago)   31m
kube-system    coredns-7db6d8ff4d-sfzs8          1/1     Running   1 (2m51s ago)   31m
kube-system    etcd-master                       1/1     Running   1 (2m51s ago)   31m
kube-system    kube-apiserver-master             1/1     Running   1 (2m51s ago)   31m
kube-system    kube-controller-manager-master    1/1     Running   1 (2m51s ago)   31m
kube-system    kube-proxy-8kd97                  1/1     Running   6 (55s ago)     7m21s
kube-system    kube-proxy-9x78m                  1/1     Running   11 (73s ago)    31m
kube-system    kube-scheduler-master             1/1     Running   1 (2m51s ago)   31m
ubuntu@master:~$
```

```
ubuntu@master:~$ kubectl get nodes
NAME      STATUS   ROLES           AGE   VERSION
master    Ready    control-plane   44m   v1.30.3
worker1   Ready    <none>          19m   v1.30.3
ubuntu@master:~$
```

Deploy service

On browser search for ngnix deployment yaml

ubuntu@master:~$ sudo nano deploy.yaml

ubuntu@master:~$ sudo cat deploy.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  selector:
    matchLabels:
      app: nginx
  replicas: 2 # tells deployment to run 2 pods matching the template
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.14.2
        ports:
        - containerPort: 80
```

ubuntu@master:~$ kubectl create -f deploy.yaml

```
      - name: nginx
        image: nginx:1.14.2
        ports:
        - containerPort: 80
ubuntu@master:~$ kubectl create -f deploy.yaml
deployment.apps/nginx-deployment created
ubuntu@master:~$ ^C
ubuntu@master:~$
```

ubuntu@master:~$ kubectl get deploy

```
ubuntu@master:~$ kubectl get deploy
NAME               READY   UP-TO-DATE   AVAILABLE   AGE
nginx-deployment   2/2     2            2           2m32s
ubuntu@master:~$ ^C
ubuntu@master:~$ ^C
ubuntu@master:~$
```

kubectl expose deployment.apps/nginix-deployment --type="LoadBalancer"

ubuntu@master:~$ kubectl get svc

```
        - containerPort: 80
ubuntu@master:~$ kubectl create -f deploy.yaml
deployment.apps/nginx-deployment created
ubuntu@master:~$ ^C
ubuntu@master:~$ kubectl get deploy
NAME               READY   UP-TO-DATE   AVAILABLE   AGE
nginx-deployment   2/2     2            2           2m32s
ubuntu@master:~$ ^C
ubuntu@master:~$ ^C
ubuntu@master:~$ kubectl expose deployment.apps/nginix-deployment --type="LoadBalance"
Error from server (NotFound): deployments.apps "nginix-deployment" not found
ubuntu@master:~$ ^C
ubuntu@master:~$ kubectl expose deployment.apps/nginix-deployment --type="LoadBalancer"
Error from server (NotFound): deployments.apps "nginix-deployment" not found
ubuntu@master:~$ kubectl expose deployment.apps/nginx-deployment --type="LoadBalancer"
service/nginx-deployment exposed
ubuntu@master:~$ kubectl get svc
NAME               TYPE           CLUSTER-IP      EXTERNAL-IP   PORT(S)        AGE
kubernetes         ClusterIP      10.96.0.1       <none>        443/TCP        57m
nginx-deployment   LoadBalancer   10.96.174.79    <pending>     80:31825/TCP   15s
ubuntu@master:~$ ^C
ubuntu@master:~$ ^C
ubuntu@master:~$
```

Go to instance,master select public ipv4

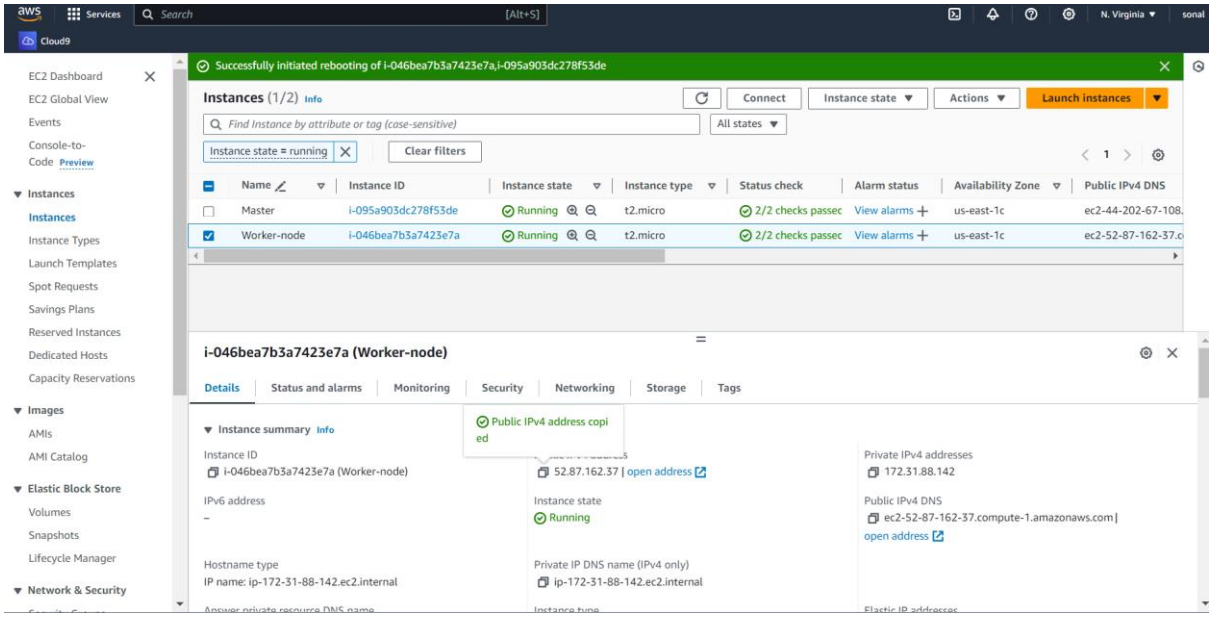

Go to brower

Ipv4:portnumber